

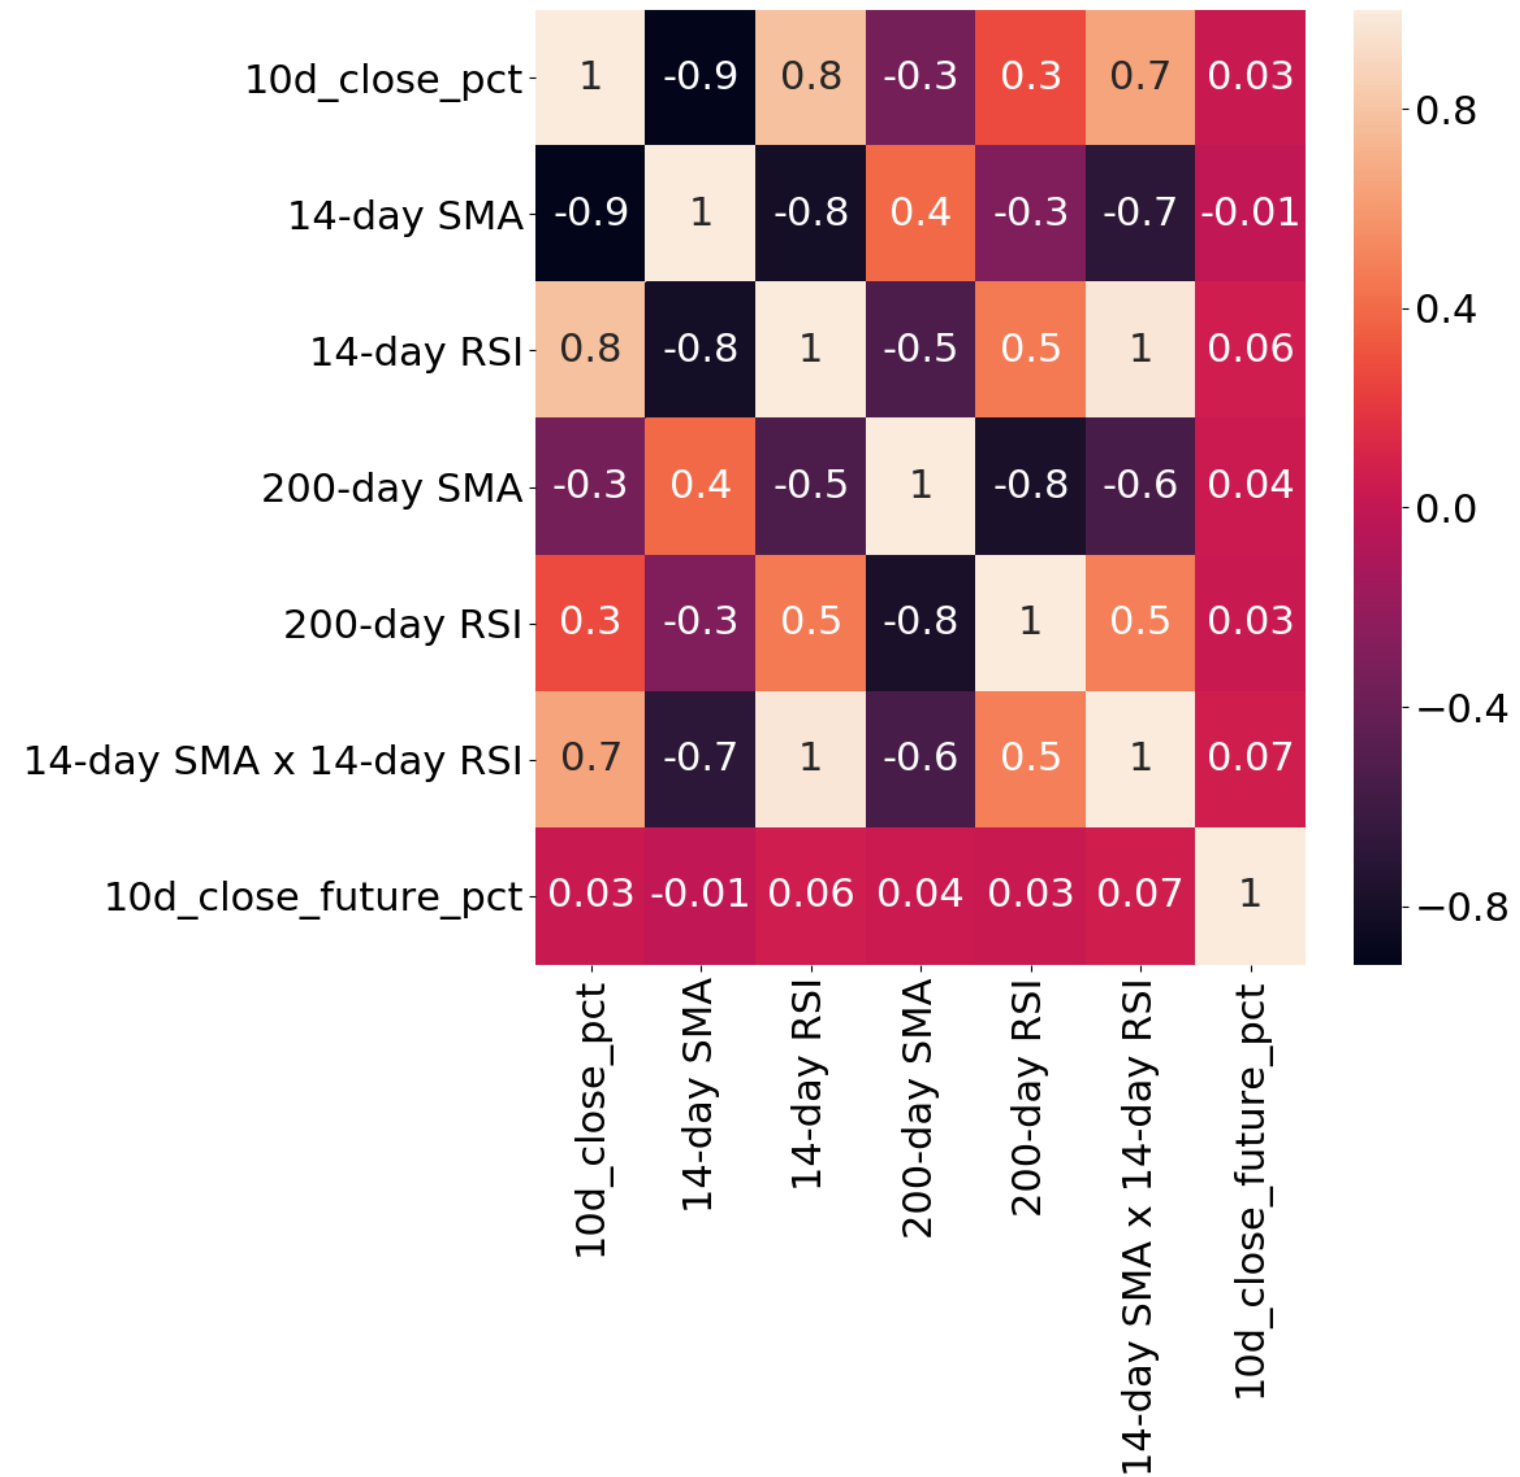
Engineering features

MACHINE LEARNING FOR FINANCE IN PYTHON



Nathan George

Data Science Professor



One problem with linear models

```
# add non-linear interaction term for a linear model  
SMAxRSI = amd_df['14-day SMA'] * amd_df['14-day RSI']
```


Some models that don't require manually creating interaction features:


Decision-tree-based models

- Random forests
- Gradient boosting

Others

- neural networks



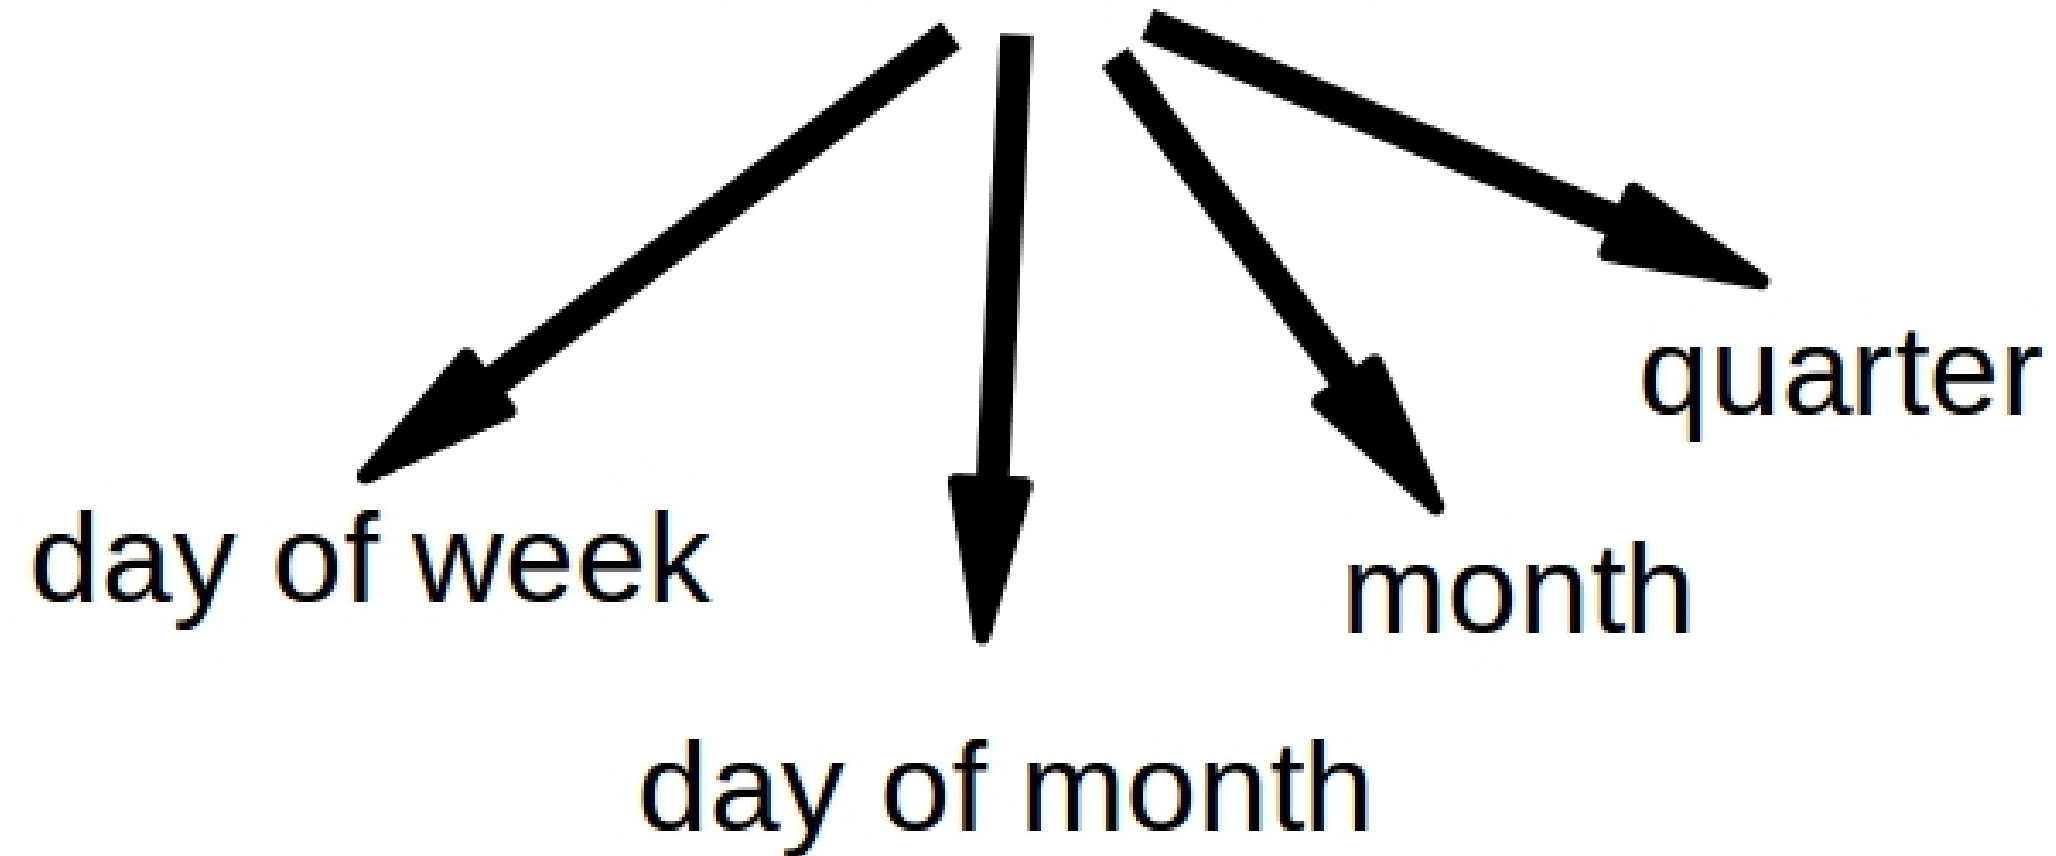


Volume features

```
amd_df['Adj_Volume_1d_change'] = amd_df['Adj_Volume'].pct_change()  
one_day_change = amd_df['Adj_Volume_1d_change'].values  
amd_df['Adj_Volume_1d_change_SMA'] = talib.SMA(one_day_change,  
                                              timeperiod=10)
```

Datetime feature engineering

2018-04-12



Extracting the day of week


```
print(amd_df.index.dayofweek)
```

```
Int64Index([2, 3, 4, 0, 1, 2, 3, 4, 0, 1,  
            ...  
            1, 2, 3, 4, 0, 1, 2, 3, 4, 0],  
            dtype='int64', name='Date', length=4807)
```


Dummies

```
days_of_week = pd.get_dummies(amd_df.index.dayofweek,  
                                prefix='weekday',  
                                drop_first=True)  
  
print(days_of_week.head())
```

	weekday_1	weekday_2	weekday_3	weekday_4
Date				
2018-04-10	1	0	0	0
2018-04-11	0	1	0	0
2018-04-12	0	0	1	0
2018-04-13	0	0	0	1
2018-04-16	0	0	0	0



Engineer some features!

MACHINE LEARNING FOR FINANCE IN PYTHON

Decision Trees


MACHINE LEARNING FOR FINANCE IN PYTHON



Nathan George


Data Science Professor

Decision trees




Decision trees

targets: [0, 0, 1, 1, 0, 1, 0]




Decision tree splits

targets: [0, 0, 1, 1, 0, 1, 0]

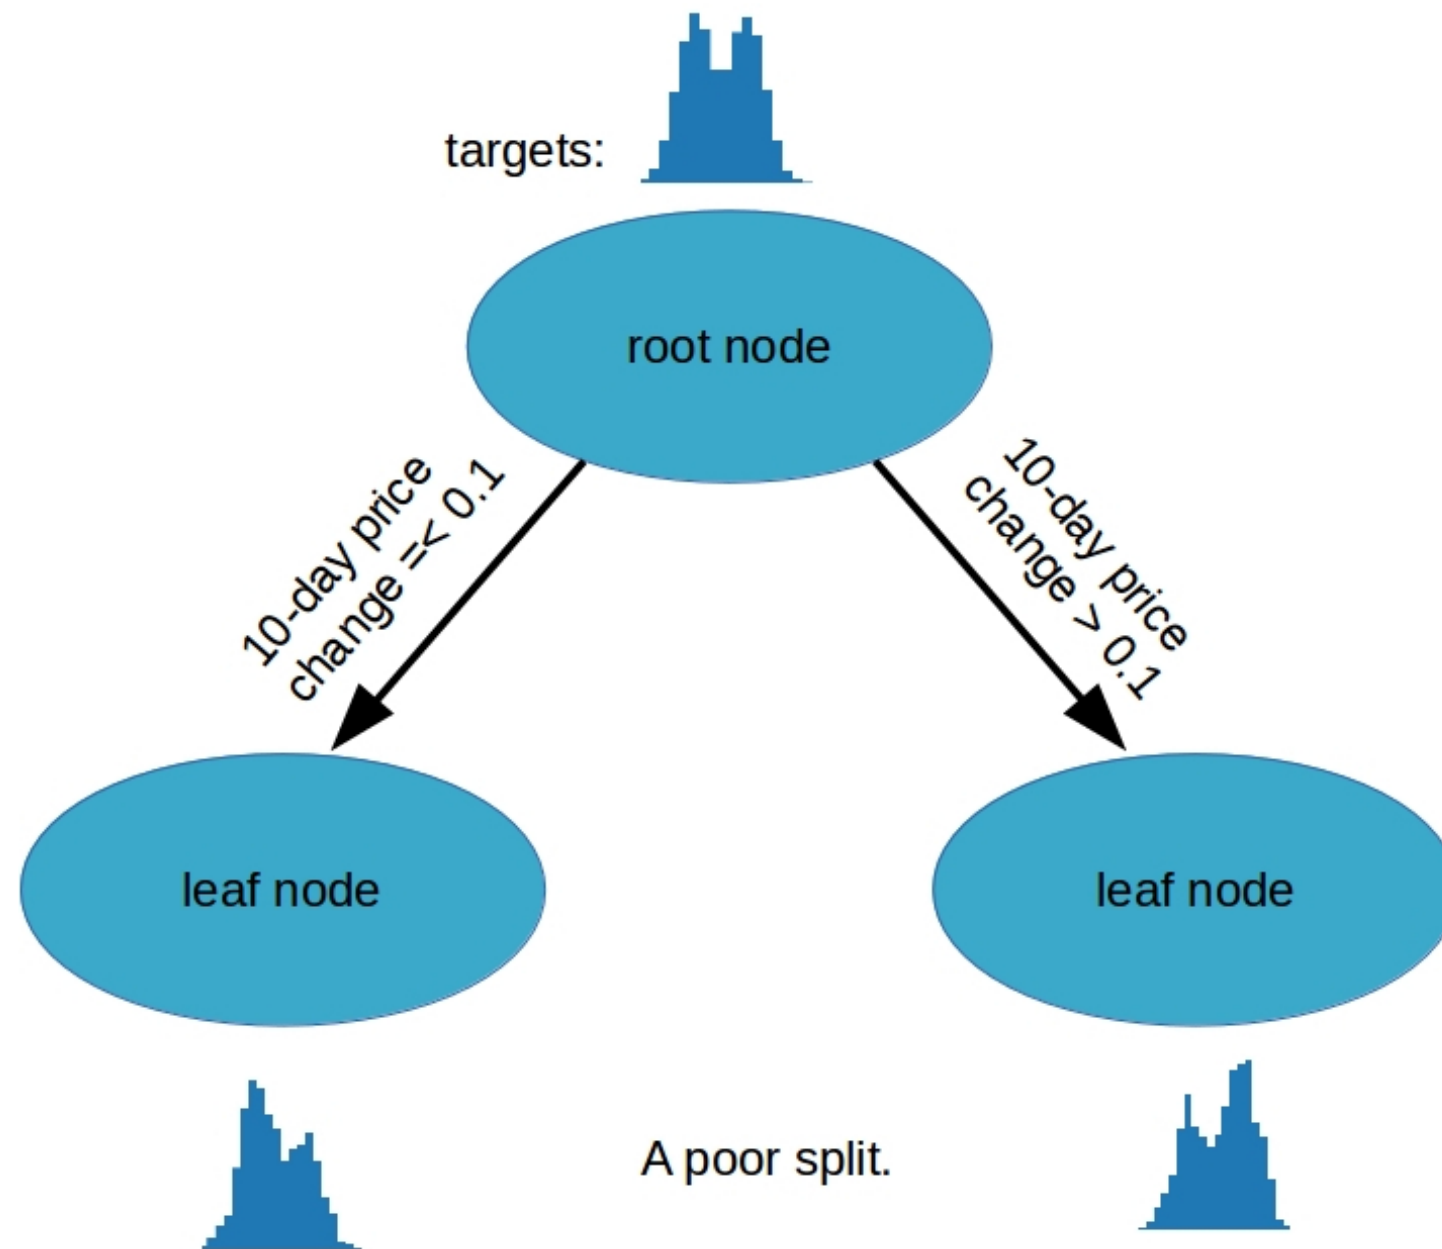


Decision tree splits

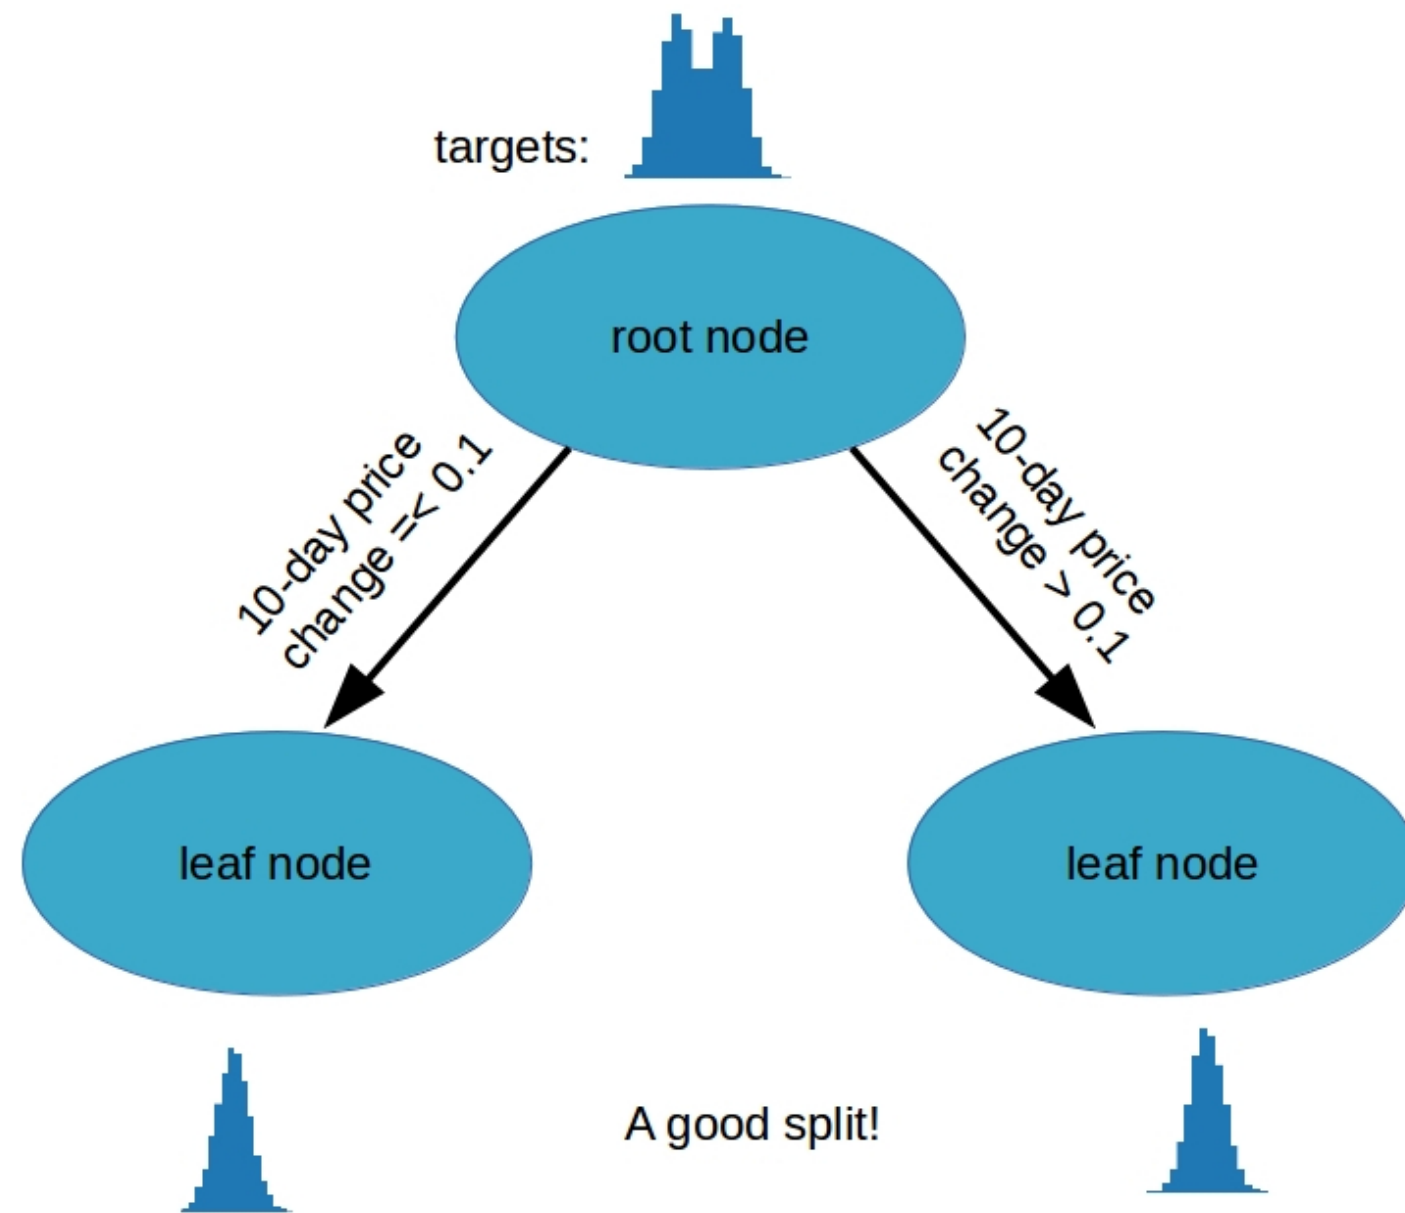
targets: [0, 0, 1, 1, 0, 1, 0]



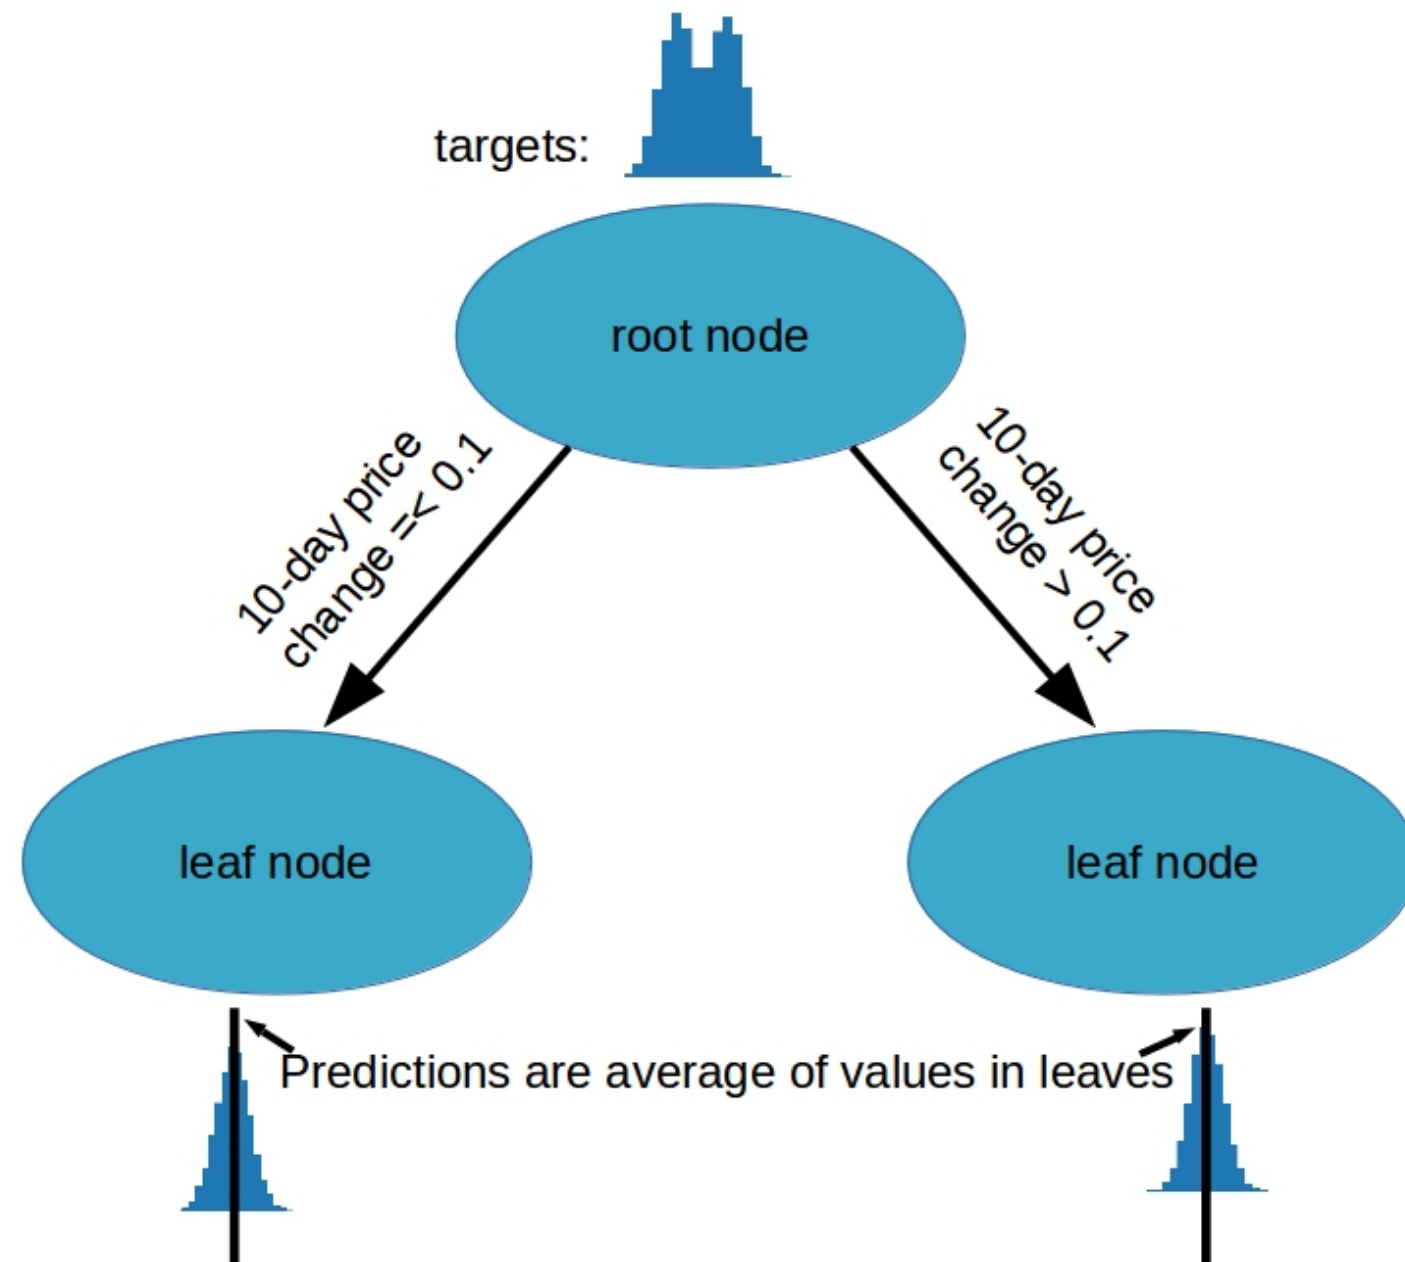
Bad tree



Good tree



Decision tree regression



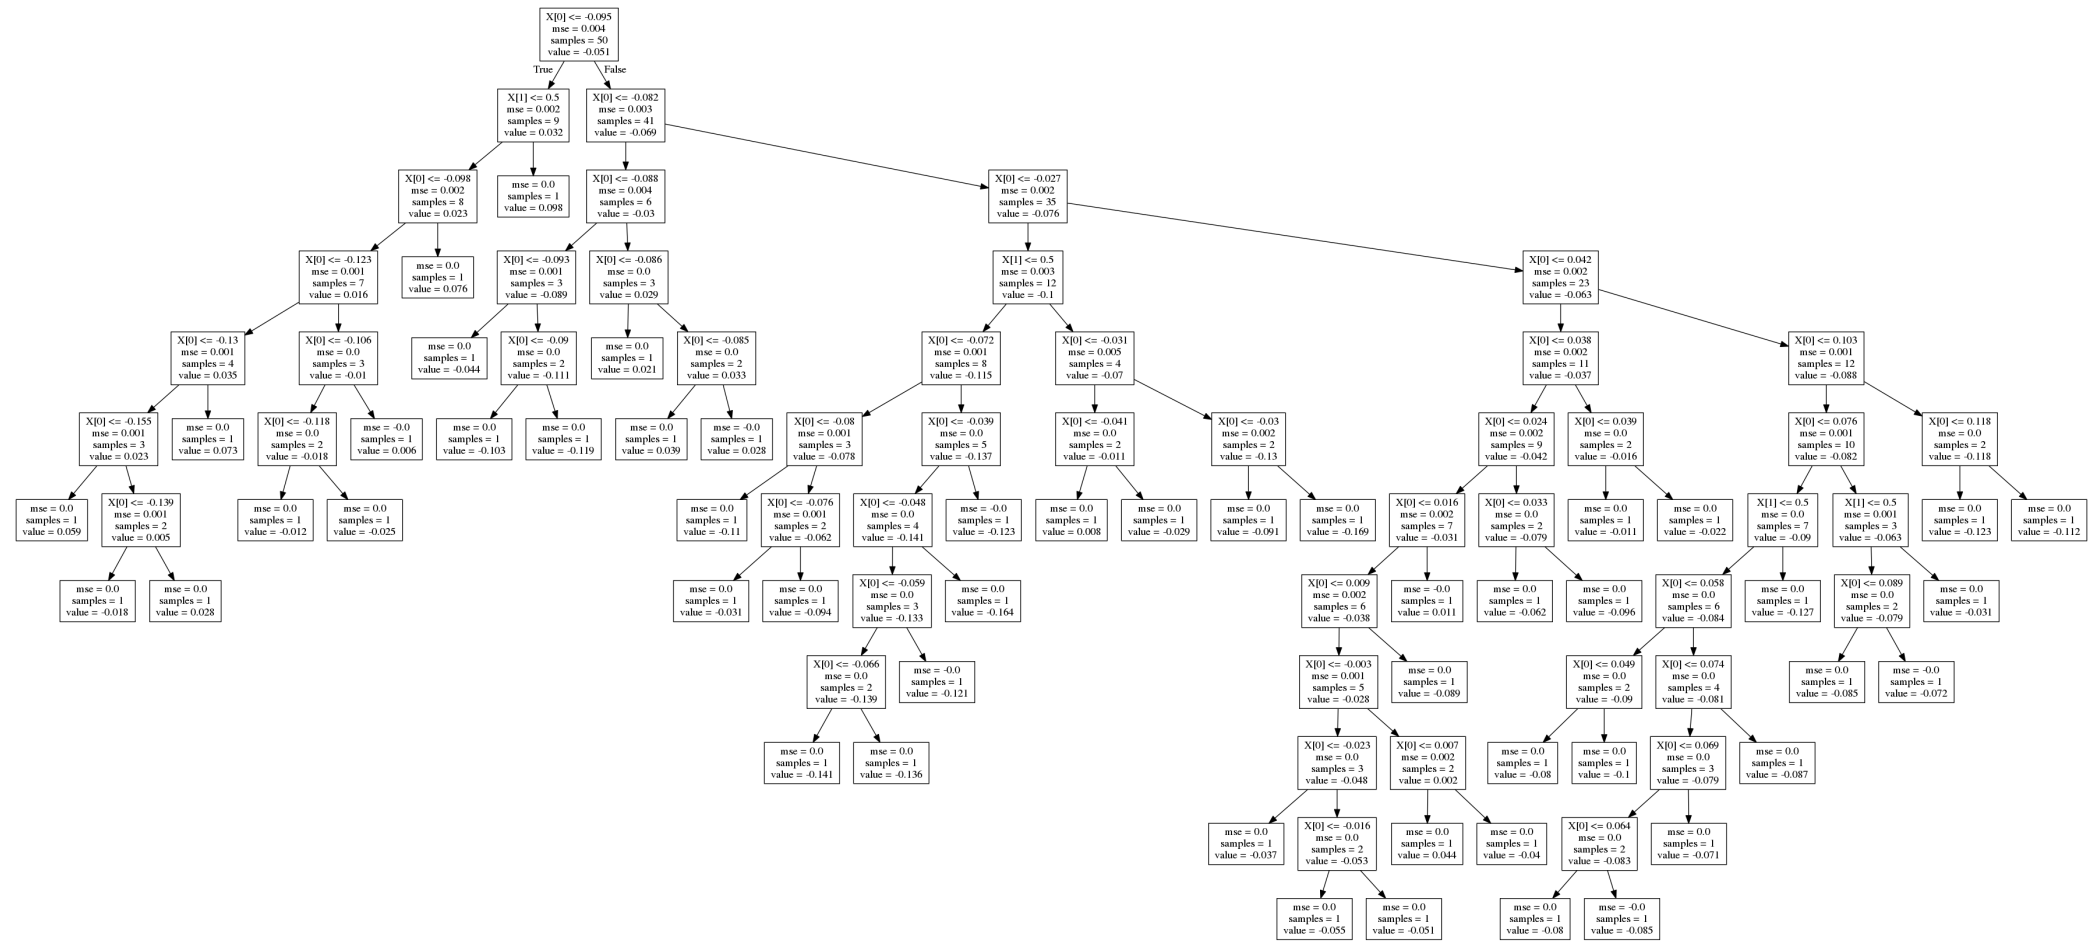
Regression trees

```
from sklearn.tree import DecisionTreeRegressor

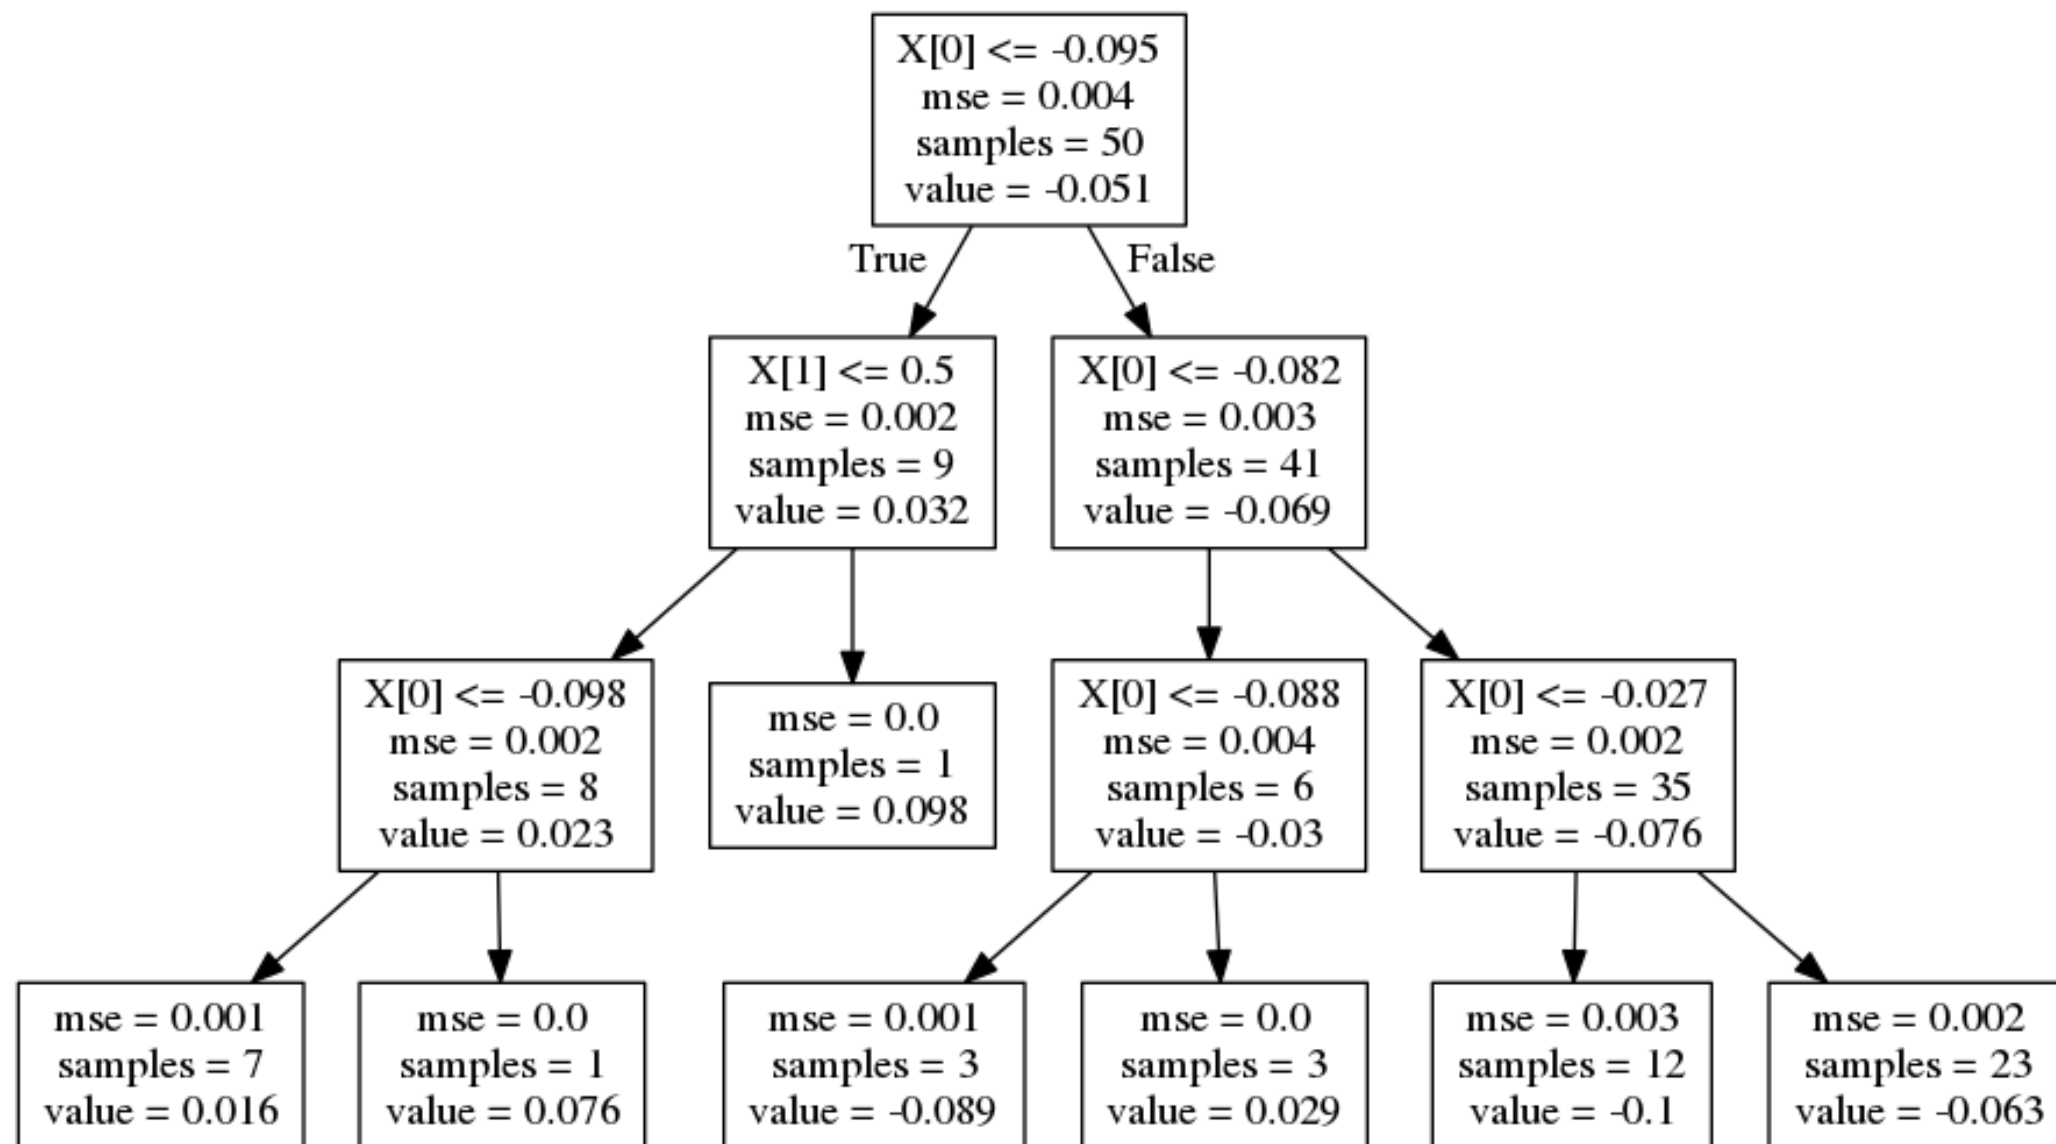
decision_tree = DecisionTreeRegressor(max_depth=5)

decision_tree.fit(train_features, train_targets)
```

Decision tree hyperparameters



Max depth of 3




Evaluate model

```
print(decision_tree.score(train_features, train_targets))  
print(decision_tree.score(test_features, test_targets))
```

```
0.6662215501032416  
-0.08917300191734268
```

```
train_predictions = decision_tree.predict(train_features)  
test_predictions = decision_tree.predict(test_features)  
plt.scatter(train_predictions, train_targets, label='train')  
plt.scatter(test_predictions, test_targets, label='test')  
plt.legend()  
plt.show()
```

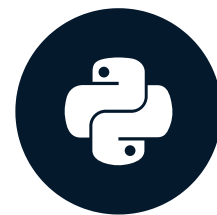


Grow some trees!

MACHINE LEARNING FOR FINANCE IN PYTHON


Random forests


MACHINE LEARNING FOR FINANCE IN PYTHON




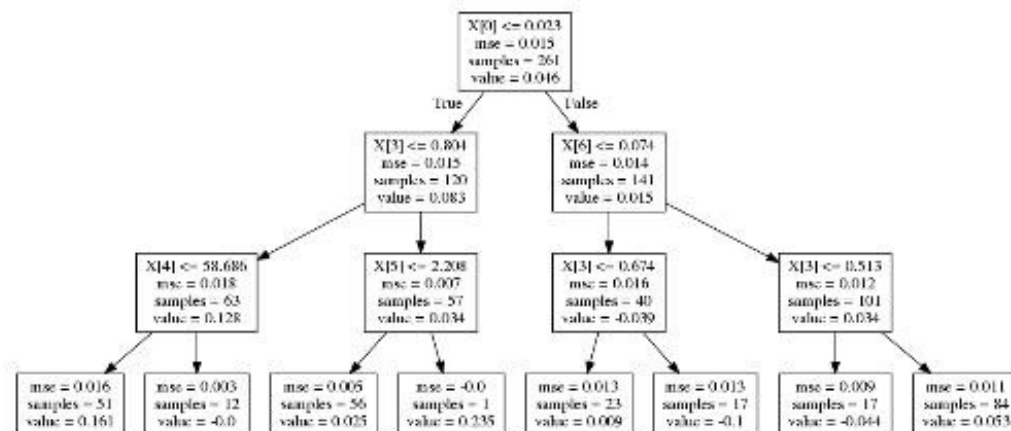


Nathan George

Data Science Professor






Random forests




Bootstrap aggregating (bagging)

training data



bootstrapped sample



Feature sampling

Random Forests

- A collection (ensemble) of decision trees
- Bootstrap aggregating (bagging)
- Sample of features at each split

sklearn implementation

```
from sklearn.ensemble import RandomForestRegressor

random_forest = RandomForestRegressor()
random_forest.fit(train_features, train_targets)
print(random_forest.score(train_features, train_targets))
```


Hyperparameters

```
random_forest = RandomForestRegressor(n_estimators=200,  
                                     max_depth=5,  
                                     max_features=4,  
                                     random_state=42)
```

ParameterGrid

```
from sklearn.model_selection import ParameterGrid

grid = {'n_estimators': [200], 'max_depth':[3, 5], 'max_feature

from pprint import pprint
pprint(list(ParameterGrid(grid)))
```

```
[{'max_depth': 3, 'max_features': 4, 'n_estimators': 200},
 {'max_depth': 3, 'max_features': 8, 'n_estimators': 200},
 {'max_depth': 5, 'max_features': 4, 'n_estimators': 200},
 {'max_depth': 5, 'max_features': 8, 'n_estimators': 200}]
```

ParameterGrid

```
test_scores = []
# loop through the parameter grid, set hyperparameters, save the scores
for g in ParameterGrid(grid):
    rfr.set_params(**g) # ** is "unpacking" the dictionary
    rfr.fit(train_features, train_targets)
    test_scores.append(rfr.score(test_features, test_targets))
# find best hyperparameters from the test score and print
best_idx = np.argmax(test_scores)
print(test_scores[best_idx])
print(ParameterGrid(grid)[best_idx])
```

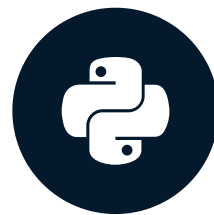
```
0.05594252725411142
{'max_depth': 5, 'max_features': 8, 'n_estimators': 200}
```

Plant some random forests!

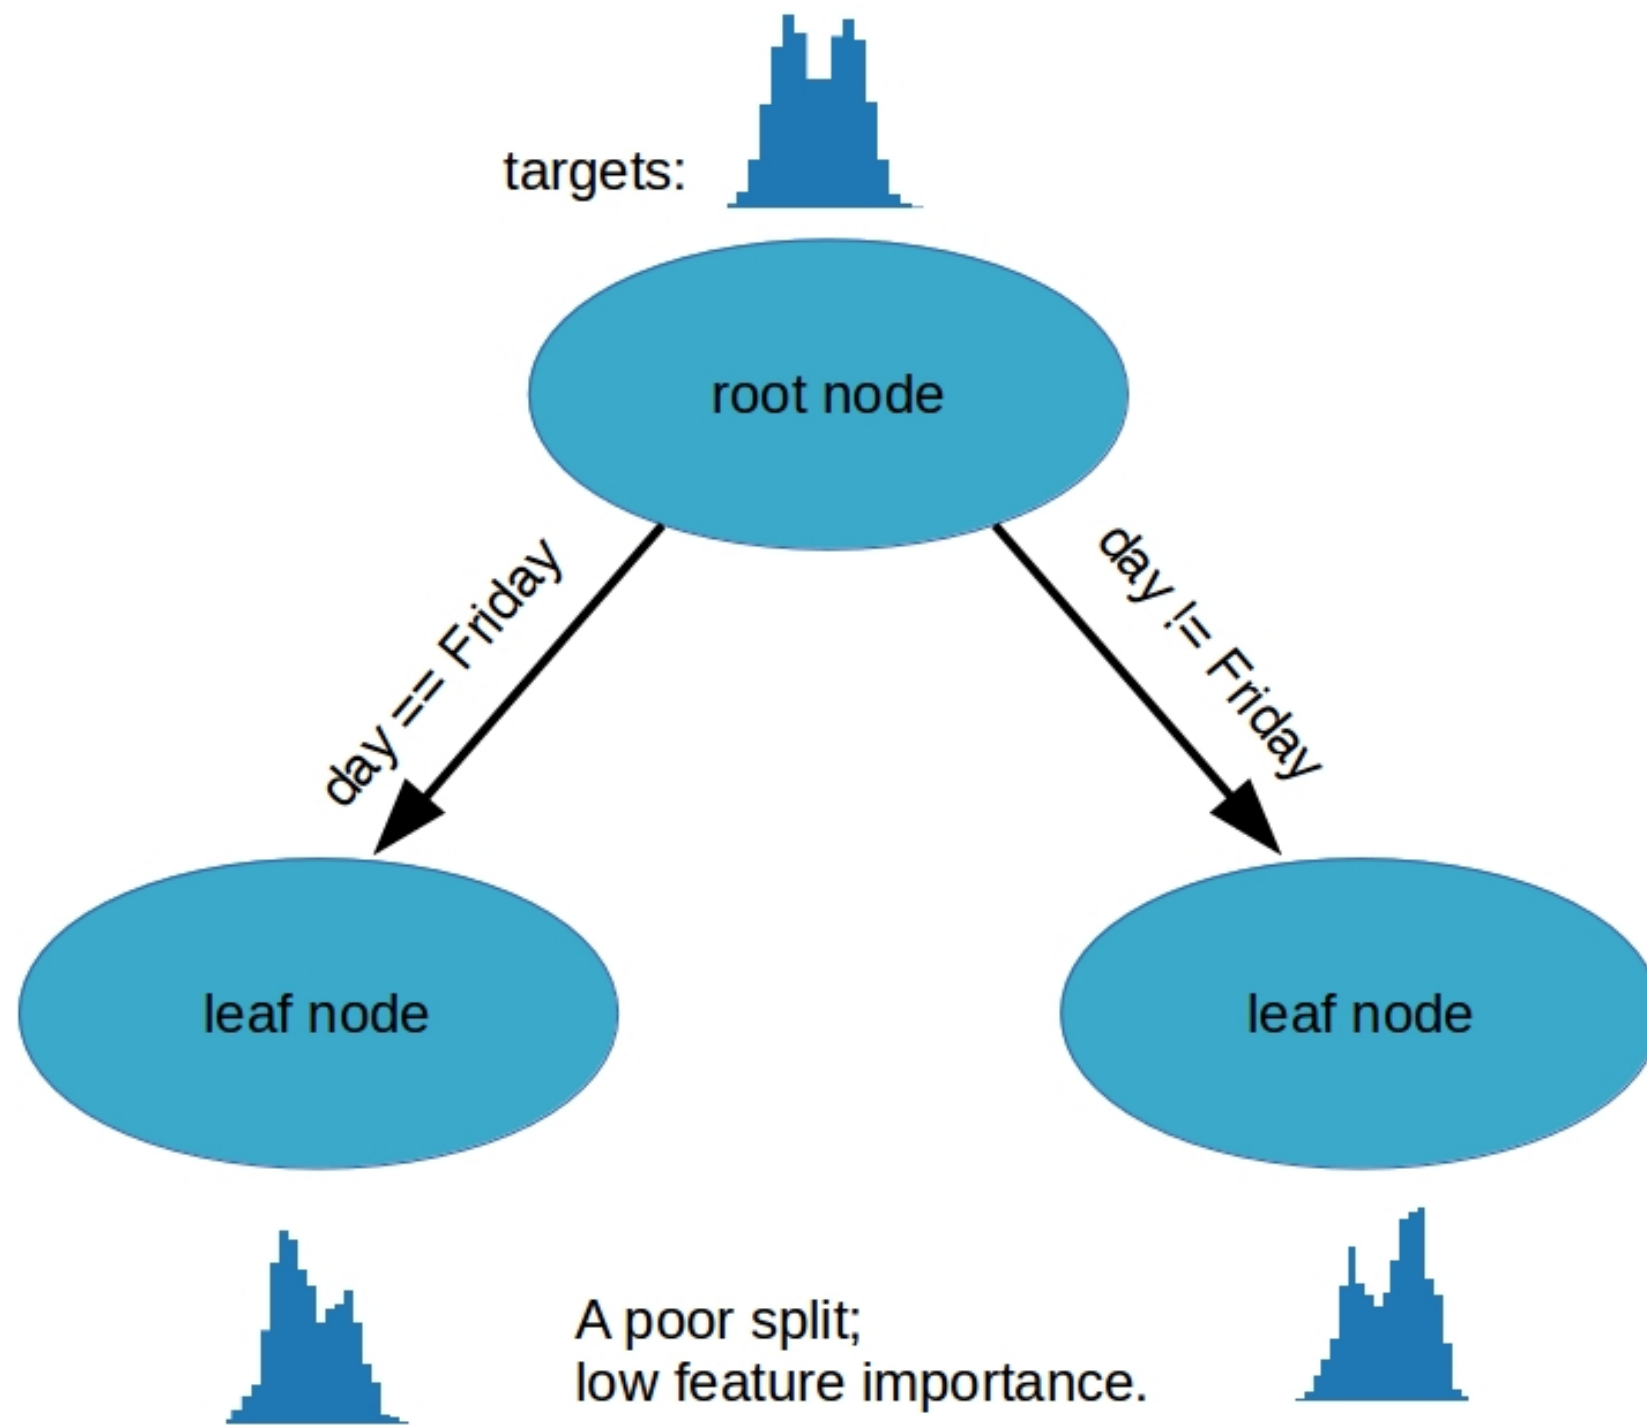
MACHINE LEARNING FOR FINANCE IN PYTHON

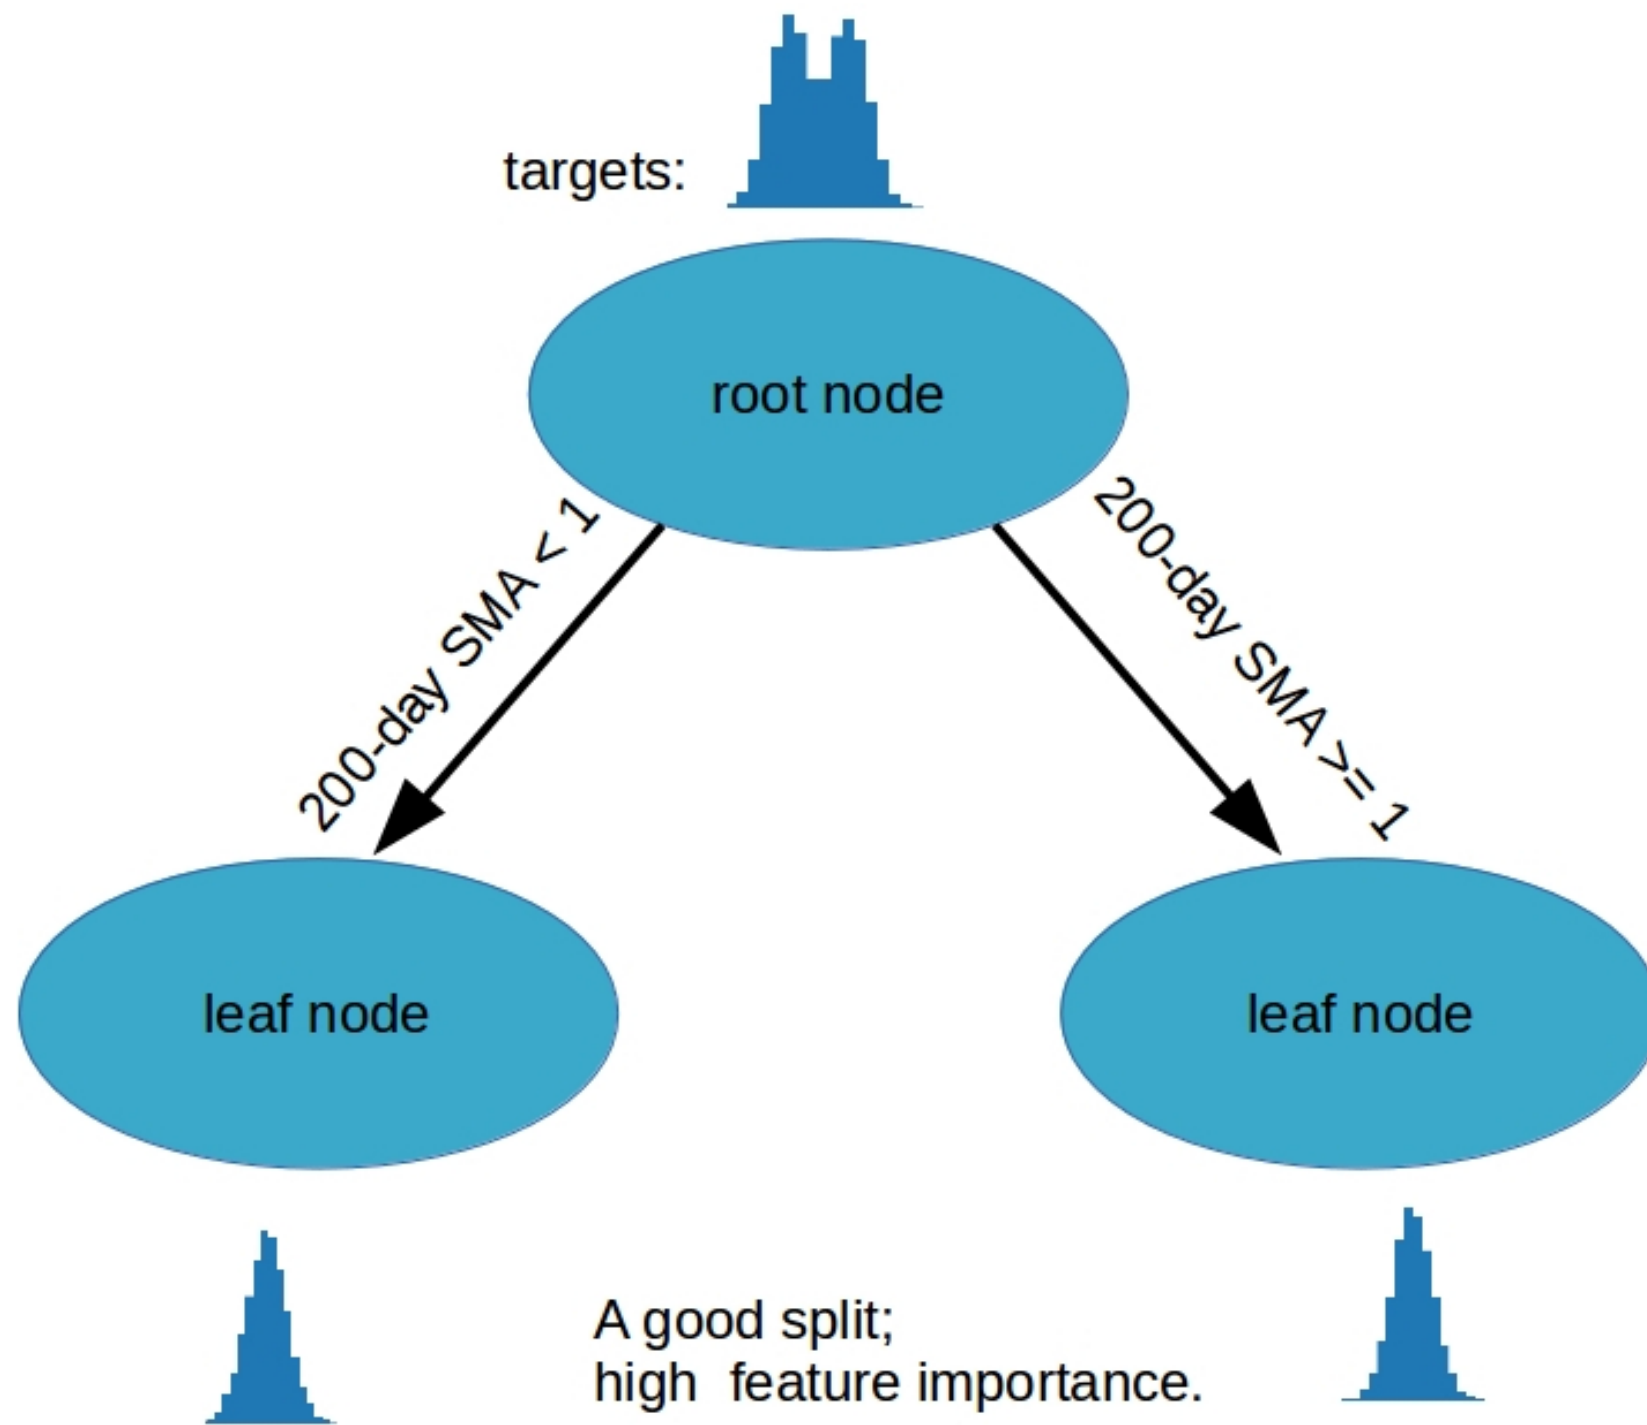
Feature importances and gradient boosting

MACHINE LEARNING FOR FINANCE IN PYTHON



Nathan George
Data Science Professor





Extracting feature importances

```
from sklearn.ensemble import RandomForestRegressor

random_forest = RandomForestRegressor()
random_forest.fit(train_features, train_targets)

feature_importances = random_forest.feature_importances_

print(feature_importances)
```

```
[0.07586547 0.10697602 0.12215955 0.23969227 0.29010304 0.0314028
 0.11977058 0.00276721 0.00246329 0.0026431  0.00615667]
```



Sorting and plotting

```
# feature importances from random forest model
importances = random_forest.feature_importances_

# index of greatest to least feature importances
sorted_index = np.argsort(importances)[::-1]
x = range(len(importances))
# create tick labels
labels = np.array(feature_names)[sorted_index]

plt.bar(x, importances[sorted_index], tick_label=labels)


# rotate tick labels to vertical
plt.xticks(rotation=90)
plt.show()
```



Linear models vs gradient boosting



<http://blog.kaggle.com/2017/01/23/a-kaggle-master-explains-gradient-boosting/>



Boosted models

Available boosted models:

- Gradient boosting
- Adaboost

Fitting a gradient boosting model

```
from sklearn.ensemble import GradientBoostingRegressor

gbr = GradientBoostingRegressor(max_features=4,
                                learning_rate=0.01,
                                n_estimators=200,
                                subsample=0.6,
                                random_state=42)

gbr.fit(train_features, train_targets)
```

Get boosted!

MACHINE LEARNING FOR FINANCE IN PYTHON