


Scaling data and KNN Regression

MACHINE LEARNING FOR FINANCE IN PYTHON



Nathan George
Data Science Professor



Feature selection: remove weekdays

```
print(feature_names)
```


```
['10d_close_pct',  
 '14-day SMA',  
 '14-day RSI',  
 '200-day SMA',  
 '200-day RSI',  
 'Adj_Volume_1d_change',  
 'Adj_Volume_1d_change_SMA',  
 'weekday_1',  
 'weekday_2',  
 'weekday_3',  
 'weekday_4']
```


```
print(feature_names[:-4])
```


```
['10d_close_pct',  
 '14-day SMA',  
 '14-day RSI',  
 '200-day SMA',  
 '200-day RSI',  
 'Adj_Volume_1d_change',  
 'Adj_Volume_1d_change_SMA']
```

Remove weekdays



```
train_features = train_features.iloc[:, :-4]  
test_features = test_features.iloc[:, :-4]
```







$$D(A, B) = \sum_i (|a_i - b_i|)^{(1/p)}$$




Scaling options


Scaling options:

- min-max
- standardization
- median-MAD
- map to arbitrary function (e.g. sigmoid, tanh)

before standardization





after standardization



sklearn's scaler

```
from sklearn.preprocessing import scaler

sc = scaler()
scaled_train_features = sc.fit_transform(train_features)
scaled_test_features = sc.transform(test_features)
```



Making subplots

```
# create figure and list containing axes
f, ax = plt.subplots(nrows=2, ncols=1)
# plot histograms of before and after scaling
train_features.iloc[:, 2].hist(ax=ax[0])
ax[1].hist(scaled_train_features[:, 2])
plt.show()
```

Scale data and use KNN!

MACHINE LEARNING FOR FINANCE IN PYTHON

Neural Networks


MACHINE LEARNING FOR FINANCE IN PYTHON



Nathan George

Data Science Professor


GPU computational power over time

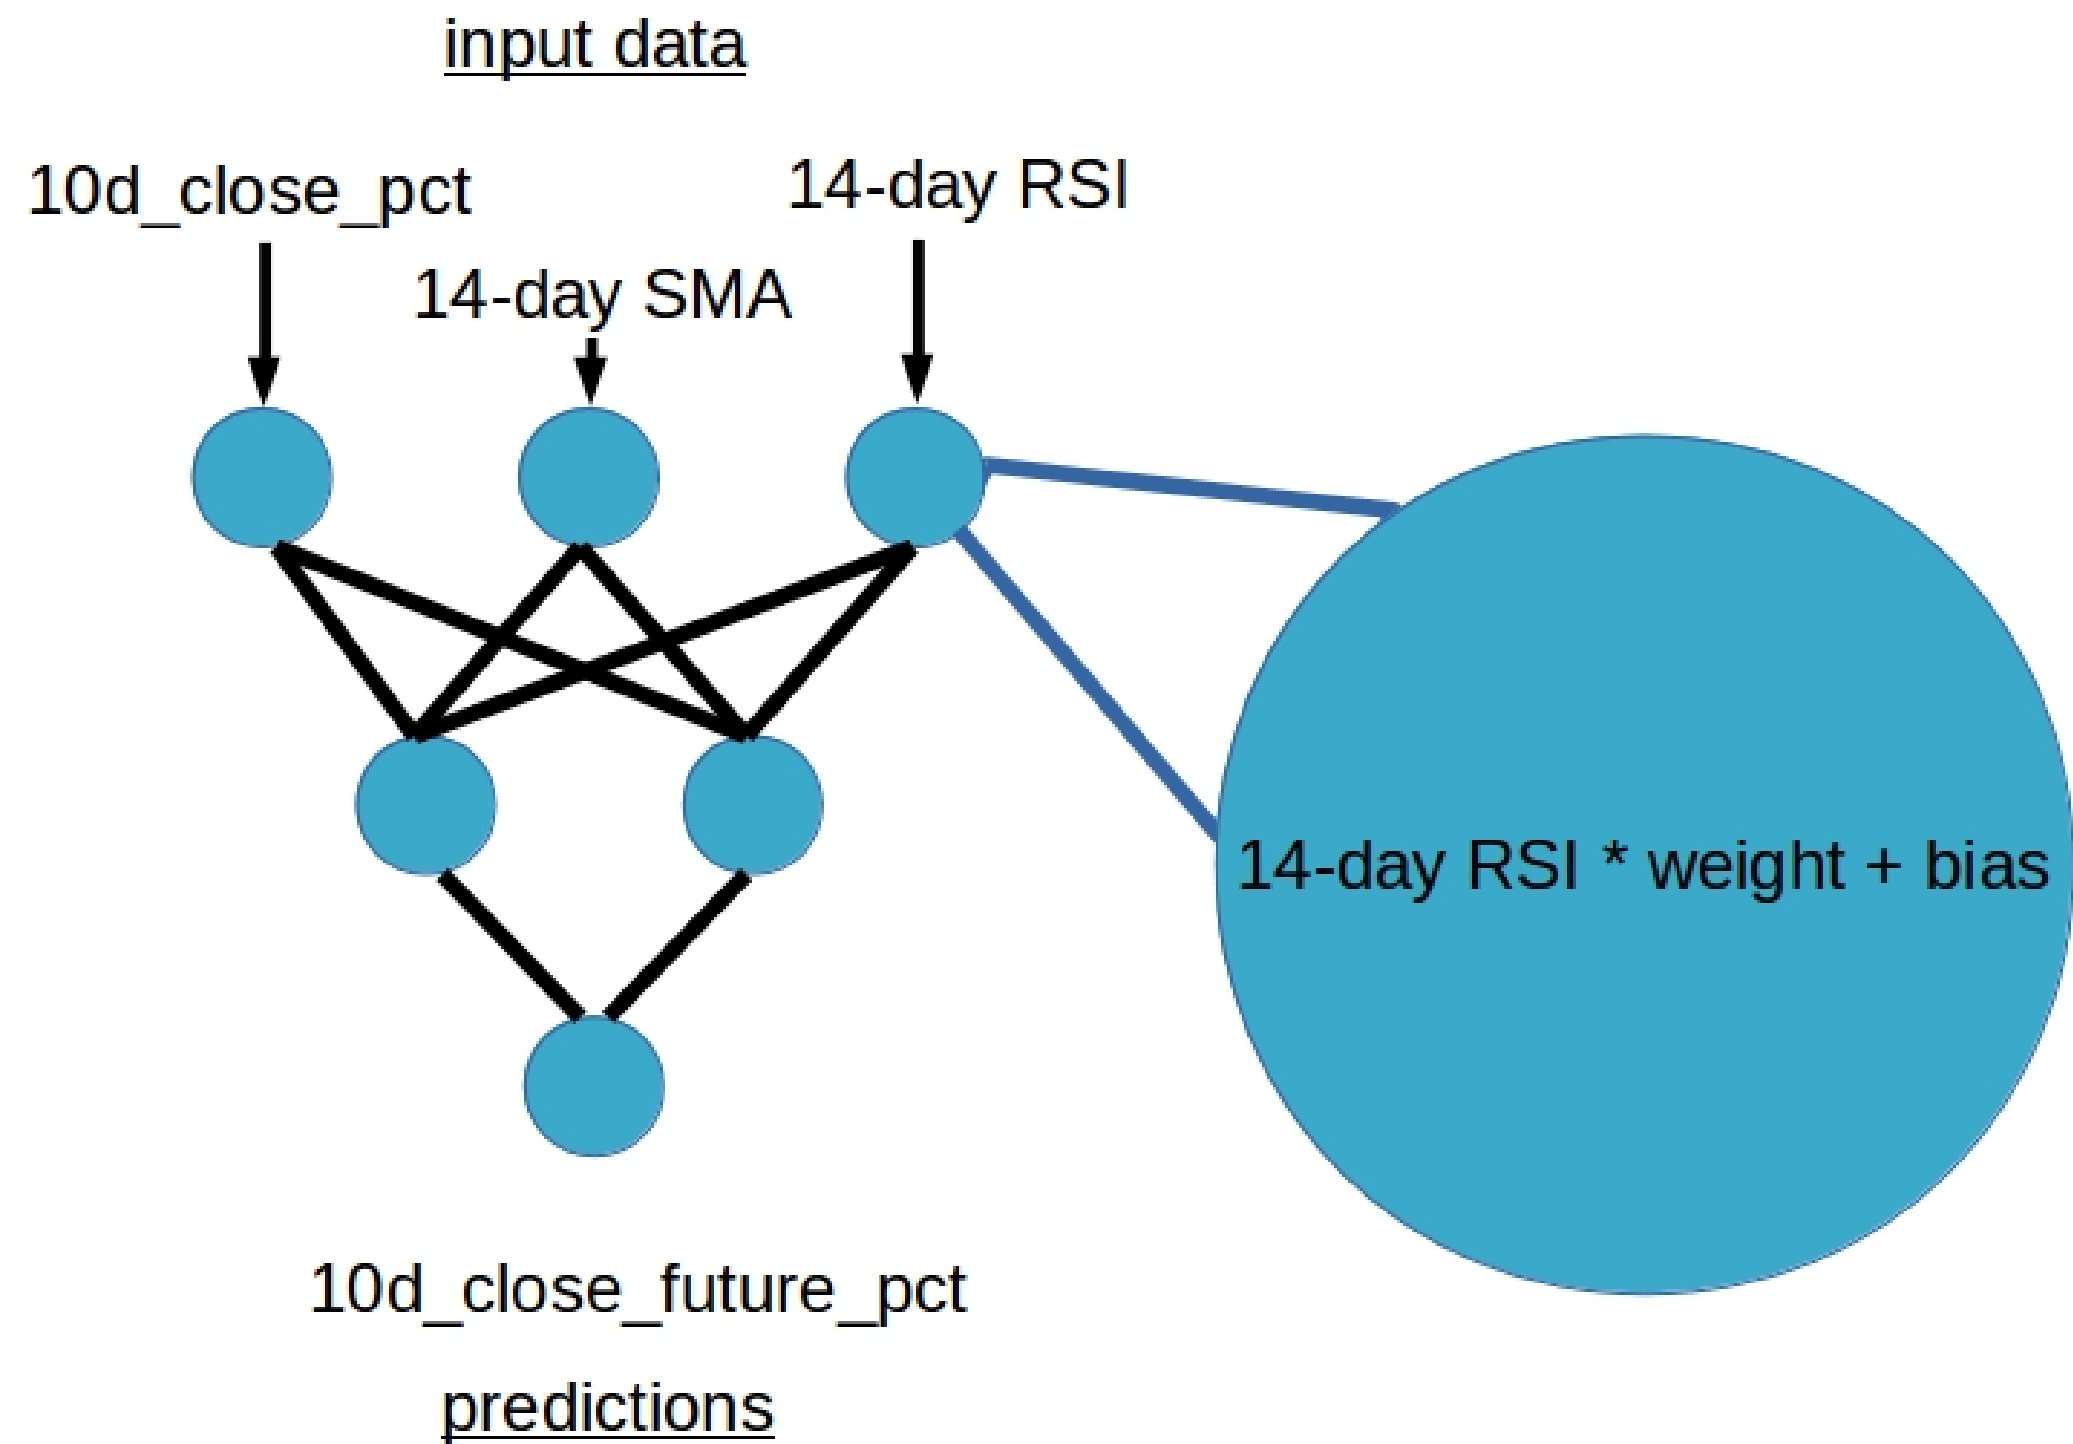


Neural networks have potential

Neural nets have:

- non-linearity
- variable interactions
- customizability





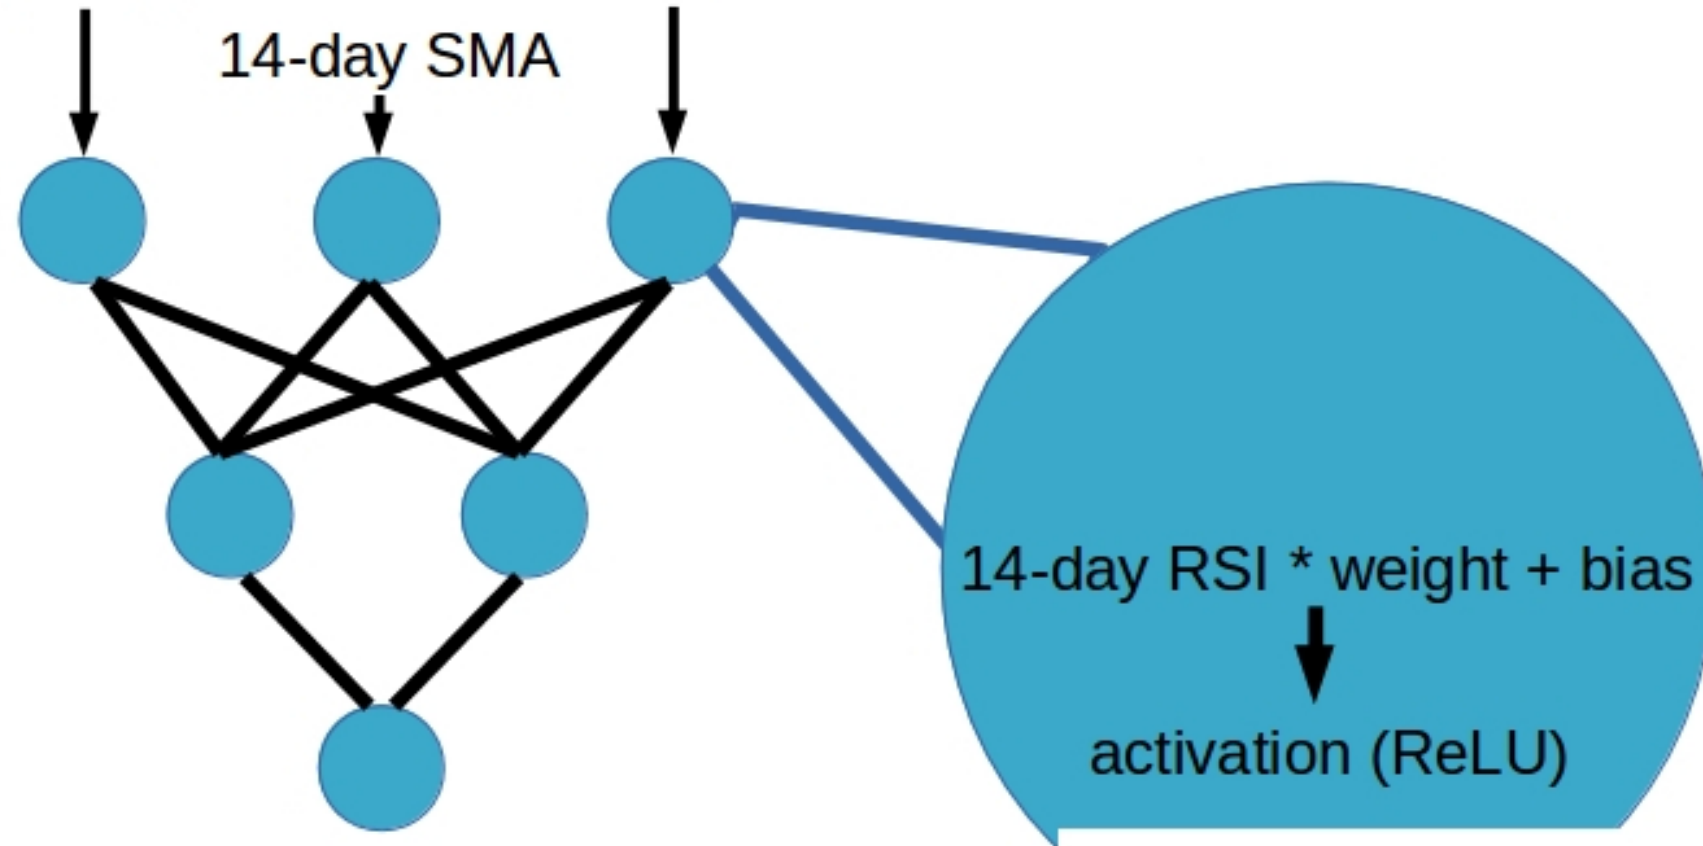
$$\sum_i w_i x_i + b$$

input data

10d_close_pct


14-day RSI


14-day SMA



10d_close_future_pct

predictions






input data

10d_close_pct

14-day RSI

14-day SMA




10d_close_future_pct



loss function (mse)

input data




input data


10d_close_pct

14-day RSI

14-day SMA



error from loss function



Keras



Implementing a neural net with keras

```
from keras.models import Sequential  
from keras.layers import Dense
```

Implementing a neural net with keras


```
from keras.models import Sequential
from keras.layers import Dense

model = Sequential()
model.add(Dense(50,
                input_dim=scaled_train_features.shape[1],
                activation='relu'))
model.add(Dense(10, activation='relu'))
model.add(Dense(1, activation='linear'))
```

Fitting the model

```
model.compile(optimizer='adam', loss='mse')  
history = model.fit(scaled_train_features,  
                    train_targets,  
                    epochs=50)
```

```
plt.plot(history.history['loss'])  
plt.title('loss:' + str(round(history.history['loss'][-1], 6)))  
plt.xlabel('epoch')  
plt.ylabel('loss')  
plt.show()
```



Checking out performance


```
from sklearn.metrics import r2_score

# calculate R^2 score
train_preds = model.predict(scaled_train_features)
print(r2_score(train_targets, train_preds))
```

```
0.4771387560719418
```


Plot performance

```
# plot predictions vs actual  
plt.scatter(train_preds, train_targets)  
plt.xlabel('predictions')  
plt.ylabel('actual')  
plt.show()
```




Make a neural net!


MACHINE LEARNING FOR FINANCE IN PYTHON

Custom loss functions

MACHINE LEARNING FOR FINANCE IN PYTHON



Nathan George
Data Science Professor



MSE with directional penalty

If prediction and target direction match:

- $\sum (y - \hat{y})^2$

If not:

- $\sum (y - \hat{y})^2 * \text{penalty}$

Implementing custom loss functions

```
import tensorflow as tf
```

Creating a function

```
import tensorflow as tf
# create loss function
def mean_squared_error(y_true, y_pred):
```

Mean squared error loss

```
import tensorflow as tf

# create loss function
def mean_squared_error(y_true, y_pred):
    loss = tf.square(y_true - y_pred)
    return tf.reduce_mean(loss, axis=-1)
```


Add custom loss to keras

```
import tensorflow as tf

# create loss function
def mean_squared_error(y_true, y_pred):
    loss = tf.square(y_true - y_pred)
    return tf.reduce_mean(loss, axis=-1)

# enable use of loss with keras
import keras.losses
keras.losses.mean_squared_error = mean_squared_error


# fit the model with our mse loss function
model.compile(optimizer='adam', loss=mean_squared_error)
history = model.fit(scaled_train_features, train_targets, epochs=50)
```

Checking for correct direction

```
tf.less(y_true * y_pred, 0)
```

Correct direction:

- $\text{neg} * \text{neg} = \text{pos}$
- $\text{pos} * \text{pos} = \text{pos}$

Wrong direction:

- $\text{neg} * \text{pos} = \text{neg}$
- $\text{pos} * \text{neg} = \text{neg}$

Using tf.where()

```
# create loss function
def sign_penalty(y_true, y_pred):
    penalty = 10.
    loss = tf.where(tf.less(y_true * y_pred, 0),
                    penalty * tf.square(y_true - y_pred),
                    tf.square(y_true - y_pred))
```

Tying it together

```
# create loss function
def sign_penalty(y_true, y_pred):
    penalty = 100.
    loss = tf.where(tf.less(y_true * y_pred, 0),
                    penalty * tf.square(y_true - y_pred),
                    tf.square(y_true - y_pred))

    return tf.reduce_mean(loss, axis=-1)
# enable use of loss with keras
keras.losses.sign_penalty = sign_penalty
```


Using the custom loss

```
# create the model
model = Sequential()
model.add(Dense(50,
                input_dim=scaled_train_features.shape[1],
                activation='relu'))
model.add(Dense(10, activation='relu'))
model.add(Dense(1, activation='linear'))
```

```
# fit the model with our custom 'sign_penalty' loss function
model.compile(optimizer='adam', loss=sign_penalty)
history = model.fit(scaled_train_features, train_targets, epochs=50)
```

The bow-tie shape

```
train_preds = model.predict(scaled_train_features)
# scatter the predictions vs actual
plt.scatter(train_preds, train_targets)
plt.xlabel('predictions')
plt.ylabel('actual')
plt.show()
```

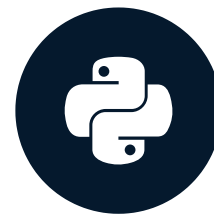


Create your own loss function!


MACHINE LEARNING FOR FINANCE IN PYTHON

Overfitting and ensembling

MACHINE LEARNING FOR FINANCE IN PYTHON




Nathan George
Data Science Professor




Simplify your model

Complex net overfits



Simpler net prevents overfitting




Neural network options

Options to combat overfitting:


- Decrease number of nodes
- Use L1/L2 regularisation
- Dropout
- Autoencoder architecture
- Early stopping
- Adding noise to data
- Max norm constraints
- Ensembling

Dropout

33% dropout



no dropout



Dropout in keras

```
from keras.layers import Dense, Dropout
model = Sequential()
model.add(Dense(500,
                input_dim=scaled_train_features.shape[1],
                activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(100, activation='relu'))
model.add(Dense(1, activation='linear'))
```

Test set comparison


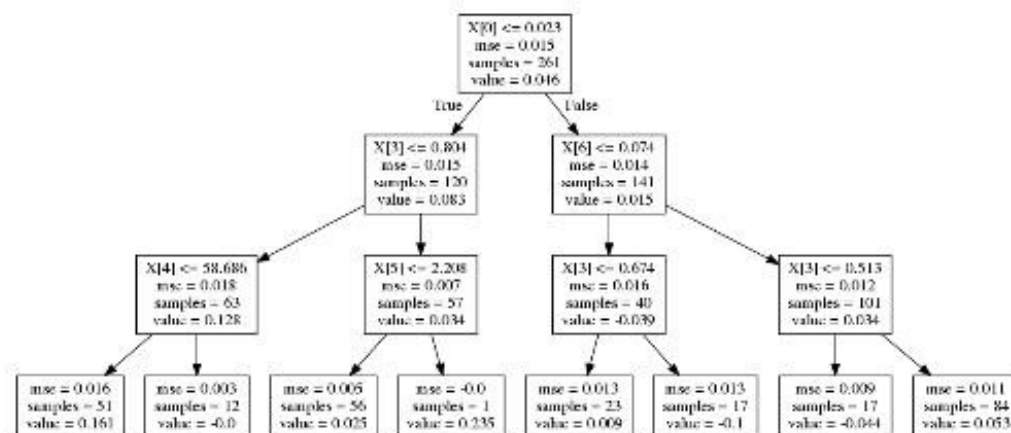

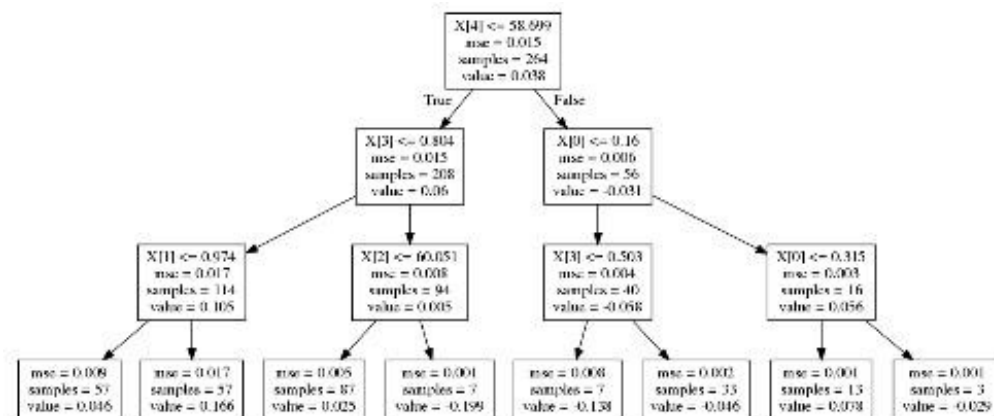
R^2 values on AMD without dropout:

- train: 0.91
- test: -0.72

With dropout:

- train: 0.46
- test: -0.22

Ensembling



Implementing ensembling

```
# make predictions from 2 neural net models
test_pred1 = model_1.predict(scaled_test_features)
test_pred2 = model_2.predict(scaled_test_features)
# horizontally stack predictions and take the average across rows
test_preds = np.mean(np.hstack((test_pred1, test_pred2)), axis=1)
```


Comparing the ensemble

Model 1 R^2 score on test set:

- -0.179

model 2:

- -0.148

ensemble (averaged predictions):

- -0.146

Dropout and ensemble!

MACHINE LEARNING FOR FINANCE IN PYTHON