

Logistic regression on sonar

MACHINE LEARNING WITH CARET IN R



Max Kuhn

Software Engineer at RStudio and
creator of caret

Classification models

- Categorical (i.e. qualitative) target variable
- Example: will a loan default?
- Still a form of supervised learning
- Use a train/test split to evaluate performance
- Use the `Sonar` dataset
- Goal: distinguish rocks from mines

Example: Sonar data

```
# Load the Sonar dataset  
library(mlbench)  
data(Sonar)
```

```
# Look at the data  
Sonar[1:6, c(1:5, 61)]
```

	V1	V2	V3	V4	V5	Class
1	0.0200	0.0371	0.0428	0.0207	0.0954	R
2	0.0453	0.0523	0.0843	0.0689	0.1183	R
3	0.0262	0.0582	0.1099	0.1083	0.0974	R
4	0.0100	0.0171	0.0623	0.0205	0.0205	R
5	0.0762	0.0666	0.0481	0.0394	0.0590	R
6	0.0286	0.0453	0.0277	0.0174	0.0384	R

Splitting the data

- Randomly split data into training and test sets
- Use a 60/40 split, instead of 80/20
- Sonar dataset is small, so 60/40 gives a larger, more reliable test set

Splitting the data

```
# Randomly order the dataset
rows <- sample(nrow(Sonar))
Sonar <- Sonar[rows, ]
```

```
# Find row to split on
split <- round(nrow(Sonar) * 0.60)
train <- Sonar[1:split, ]
test <- Sonar[(split + 1):nrow(Sonar), ]
```

```
# Confirm test set size
nrow(train) / nrow(Sonar)
```

```
0.6009615
```

Let's practice!

MACHINE LEARNING WITH CARET IN R

Confusion matrix

MACHINE LEARNING WITH CARET IN R



Zach Mayer

Data Scientist at DataRobot and co-author of caret

Confusion matrix

		Reference	
Prediction		Yes	No
	Yes	True positive	False positive
	No	False negative	True Negative

Confusion matrix

```
# Fit a model
model <- glm(Class ~ ., family = binomial(link = "logit"), train)
p <- predict(model, test, type = "response")
summary(p)
```

```
   Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.0000  0.0000  0.9885  0.5296  1.0000  1.0000
```

```
# Turn probabilities into classes and look at their frequencies
p_class <- ifelse(p > 0.50, "M", "R")
table(p_class)
p_class
```

```
 M   R
44  39
```

Confusion matrix

- Make a 2-way frequency table
- Compare predicted vs. actual classes

```
# Make simple 2-way frequency table  
table(p_class, test[["Class"]])
```

```
p_class  M  R  
      M 13 31  
      R 30  9
```

Confusion matrix

```
# Use caret's helper function to calculate additional statistics
confusionMatrix(p_class, test[["Class"]])
```

```
      Reference
Prediction M  R
M 13 31
R 30  9

      Accuracy : 0.2651
      95% CI : (0.1742, 0.3734)
No Information Rate : 0.5181
P-Value [Acc > NIR] : 1

      Kappa : -0.4731
McNemar's Test P-Value : 1

      Sensitivity : 0.3023
      Specificity : 0.2250
Pos Pred Value : 0.2955
Neg Pred Value : 0.2308
```

Let's practice!

MACHINE LEARNING WITH CARET IN R

Class probabilities and predictions

MACHINE LEARNING WITH CARET IN R



Zach Mayer

Data Scientist at DataRobot and co-
author of caret

Different thresholds

- Not limited to 50% threshold
 - 10% would catch more mines with less certainty
 - 90% would catch fewer mines with more certainty
- Balance true positive and false positive rates
- Cost-benefit analysis

Confusion matrix

```
# Use a larger cutoff
p_class <- ifelse(p > 0.99, "M", "R")
table(p_class)
```

```
p_class
 M  R
41 42
```

```
# Make simple 2-way frequency table
table(p_class, test[["Class"]])
```

```
p_class  M  R
      M 13 28
      R 30 12
```

Confusion matrix with caret

```
# Use caret to produce confusion matrix
confusionMatrix(p_class, test[["Class"]])
```

```
      Reference
Prediction M  R
M 13 28
R 30 12

      Accuracy : 0.3012
      95% CI : (0.2053, 0.4118)
No Information Rate : 0.5181
P-Value [Acc > NIR] : 1.0000

      Kappa : -0.397
McNemar's Test P-Value : 0.8955

      Sensitivity : 0.3023
      Specificity : 0.3000
Pos Pred Value : 0.3171
Neg Pred Value : 0.2857
```


Let's practice!

MACHINE LEARNING WITH CARET IN R

Introducing the ROC curve

MACHINE LEARNING WITH CARET IN R



Zach Mayer

Data Scientist at DataRobot and co-
author of caret

The challenge


- Many possible classification thresholds
- Requires manual work to choose
- Easy to overlook a particular threshold
- Need a more systematic approach

ROC curves

- Plot true/false positive rate at every possible threshold
- Visualize tradeoffs between two extremes (100% true positive vs. 0% false positive)
- Result is an ROC curve
- Developed as a method for analyzing radar signals

An example ROC curve

```
# Create ROC curve  
library(caTools)  
colAUC(p, test[["Class"]], plotROC = TRUE)
```



Let's practice!

MACHINE LEARNING WITH CARET IN R

Area under the curve (AUC)


MACHINE LEARNING WITH CARET IN R



Zach Mayer

Data Scientist at DataRobot and co-author of caret

From ROC to AUC



Defining AUC

- Single-number summary of model accuracy
- Summarizes performance across all thresholds
- Rank different models within the same dataset

Defining AUC

- Ranges from 0 to 1
 - 0.5 = random guessing
 - 1 = model always right
 - 0 = model always wrong
- Rule of thumb: AUC as a letter grade
 - 0.9 = "A"
 - 0.8 = "B"
 - ...

Let's practice!

MACHINE LEARNING WITH CARET IN R